

N-grams in practice: spelling correction

- How would a spell checker know that there is a spelling error in the following sentence?
I like pizza with spinach
- Or this one?
Zoo animals on the lose

We want:

$P(I \text{ like pizza with spinach}) > P(I \text{ like pizza wit spinach})$
 $P(\text{Zoo animals on the loose}) > P(\text{Zoo animals on the lose})$

More applications for language models

- Spelling correction
- Speech recognition
- Machine translation
- Predictive text
- Text recognition (OCR, handwritten)
- Information retrieval
- Question answering
- Text classification
- In general, pre-trained (neural) language models can bring additional linguistic/world knowledge to almost any NLP task

Assigning probabilities to sentences

count and divide?

How do we calculate the probability of a sentence like
 $P(I \text{ like pizza with spinach})$

- Can we count the occurrences of the sentence, and divide it by the total number of sentences (in a large corpus)?
- Short answer: No.
 - Many sentences are not observed even in very large corpora
 - For the ones observed in a corpus, probabilities will not reflect our intuitions, or will not be useful in most applications



Example: applying the chain rule

$$P(I \text{ like pizza with spinach}) = P(\text{like} | I) \times P(\text{pizza} | I \text{ like}) \times P(\text{with} | I \text{ like pizza}) \times P(\text{spinach} | I \text{ like pizza with})$$

- Did we solve the problem?
- Not really, the last term is equally difficult to estimate

Maximum-likelihood estimation (MLE)

- The MLE of n-gram probabilities is based on their frequencies in a corpus
- We are interested in conditional probabilities of the form: $P(w_i | w_1, \dots, w_{i-1})$, which we estimate using

$$P(w_i | w_1, \dots, w_{i-1}) = \frac{C(w_1, \dots, w_{i-1}, w_i)}{C(w_1, \dots, w_{i-1})}$$

where, $C()$ is the frequency (count) of the sequence in the corpus.

- For example, the probability $P(\text{like} | I)$ would be

$$P(\text{like} | I) = \frac{C(I \text{ like})}{C(I)}$$

= $\frac{\text{number of times I like occurs in the corpus}}{\text{number of times I occurs in the corpus}}$

- A **language model** answers the question *how likely is a sequence of words in a given language?*
- They assign scores, typically probabilities, to sequences (of words, letters, ...)
- n-gram language models** are the 'classical' approach to language modeling
- The main idea is to estimate probabilities of sequences, using the probabilities of words given a limited history
- As a bonus we get the answer for *what is the most likely word given previous words?*

N-grams in practice: speech recognition



We want:

$P(\text{recognize speech}) > P(\text{wreck a nice beach})$

* Reproduced from (Billock, 1999)

Our aim

We want to solve two related problems:

- Given a sequence of words $w = (w_1 w_2 \dots w_m)$, what is the probability of the sequence $P(w)$?
 (machine translation, automatic speech recognition, spelling correction)
- Given a sequence of words $w_1 w_2 \dots w_{m-1}$, what is the probability of the next word $P(w_m | w_1 \dots w_{m-1})$?
 (predictive text)

Assigning probabilities to sentences

applying the chain rule

- The solution is to **decompose**
 We use probabilities of parts of the sentence (words) to calculate the probability of the whole sentence
- Using the chain rule of probability (without loss of generality), we can write

$$P(w_1, w_2, \dots, w_m) = P(w_2 | w_1) \times P(w_3 | w_1, w_2) \times \dots \times P(w_m | w_1, w_2, \dots, w_{m-1})$$

Example: bigram probabilities of a sentence

with first-order Markov assumption

$$P(I \text{ like pizza with spinach}) = P(\text{like} | I) \times P(\text{pizza} | \text{like}) \times P(\text{with} | \text{pizza}) \times P(\text{spinach} | \text{with})$$

- Now, hopefully, we can count them in a corpus

MLE estimation of an n-gram language model

An n-gram model conditioned on $n-1$ previous words.

$$\begin{aligned} \text{unigram} \quad P(w_1) &= \frac{C(w_1)}{N} \\ \text{bigram} \quad P(w_i) &= P(w_i | w_{i-1}) = \frac{C(w_{i-1} w_i)}{C(w_{i-1})} \\ \text{trigram} \quad P(w_i) &= P(w_i | w_{i-2} w_{i-1}) = \frac{C(w_{i-2} w_{i-1} w_i)}{C(w_{i-2} w_{i-1})} \end{aligned}$$

Parameters of an n-gram model are these conditional probabilities.

Unigrams

Unigrams are simply the single words (or tokens).

A small corpus

I 'm sorry . Dave .
I 'm afraid I can 't do that .

When tokenized, we have 15 tokens, and 11 types.

Unigram counts

I	3	,	1	afraid	1	do	1
'm	2	Dave	1	can	1	that	1
sorry	1	.	2	't	1		

Traditionally, can't is tokenized as can't (similar to have'n, a_n't etc.), but for our purposes, can't is more readable.

Unigram probability of a sentence

Unigram counts

I	3	,	1	afraid	1	do	1
'm	2	Dave	1	can	1	that	1
sorry	1	.	2	't	1		

$$P(I \text{ 'm sorry . Dave .})$$

$$= P(I) \times P('m) \times P(sorry) \times P(.) \times P(Dave) \times P(.)$$

$$= \frac{1}{15} \times \frac{2}{15} \times \frac{1}{15} \times \frac{1}{15} \times \frac{1}{15} \times \frac{1}{15}$$

$$= 0.00000105$$

- $P(I \text{ 'm I . sorry Dave}) = ?$
- Where did all the probability mass go?
- What is the most likely sentence according to this model?

N-gram models define probability distributions

- An n-gram model defines a probability distribution over words

$$\sum_{w \in V} P(w) = 1$$

- They also define probability distributions over word sequences of equal size. For example (length 2),


$$\sum_{w_1 \in V} \sum_{w_2 \in V} P(w_1)P(w_2) = 1$$

- What about sentences?


word	prob
I	0.200
'm	0.133
,	0.133
't	0.067
Dave	0.067
afraid	0.067
can	0.067
do	0.067
sorry	0.067
that	0.067
	1.000

Unigram probabilities in a (slightly) larger corpus

MLE probabilities in the Universal Declaration of Human Rights



Unigram probabilities



Bigrams

Bigrams are overlapping sequences of two tokens.



Bigram counts

ngram	freq	ngram	freq	ngram	freq
I 'm	2	, Dave	1	't do	1
'm sorry	1	Dave .	1	I can	1
sorry .	1	'm afraid	1	can 't	1
				that .	1

- What about the bigram ' . I '?

Sentence boundary markers

If we want sentence probabilities, we need to mark them.

(s) I 'm sorry , Dave . (/s)
(s) I 'm afraid I can 't do that . (/s)

- The bigram ' (s) ' is not the same as the unigram ' I '
- Including (/s) allows us to predict likely words at the beginning of a sentence
- Including (/s) allows us to assign a proper probability distribution to sentences

Calculating bigram probabilities

recap with some more detail

We want to calculate $P(w_2 | w_1)$. From the chain rule:

$$P(w_2 | w_1) = \frac{P(w_1, w_2)}{P(w_1)}$$

and, the MLE


$$P(w_2 | w_1) = \frac{C(w_1, w_2)}{C(w_1)} = \frac{C(w_1, w_2)}{C(w_1)}$$

$P(w_2 | w_1)$ is the probability of w_2 given the previous word is w_1

$P(w_1, w_2)$ is the probability of the sequence w_1, w_2

$P(w_1)$ is the probability of w_1 occurring as the first item in a bigram, not its unigram probability

Sentence probability: bigram vs. unigram



$$P_{uni}((s) \text{ I 'm sorry , Dave . (/s)}) = 2.83 \times 10^{-9}$$

$$P_{bi}((s) \text{ I 'm sorry , Dave . (/s)}) = 0.33$$

Unigram vs. bigram probabilities


in sentences and non-sentences

w	I	'm	sorry	,	Dave	.
P_{uni}	0.18	0.12	0.06	0.06	0.06	0.12
P_{bi}	1.00	0.67	0.50	1.00	1.00	0.33

w	,	'm	I	.	sorry	Dave
P_{uni}	0.06	0.12	0.18	0.12	0.06	0.06
P_{bi}	0.00	0.00	0.00	0.00	0.00	0.00

w	I	'm	afraid	,	Dave	.
P_{uni}	0.18	0.12	0.06	0.06	0.06	0.12
P_{bi}	1.00	0.67	0.50	0.00	1.00	1.00

Bigram models as weighted finite-state automata



Trigrams


(s) (s) I 'm sorry, Dave. (/s)
(s) (s) I 'm afraid I can't do that. (/s)

Trigram counts

ngram	freq	ngram	freq	ngram	freq
(s) (s) I	2	do that.	1	that. (/s)	1
(s) I 'm	2	I 'm sorry	1	'm sorry, /	1
sorry, Dave	1	, Dave.	1	Dave. (/s)	1
I 'm afraid	1	'm afraid I	1	afraid I can	1
I can't	1	can't do	1	't do that	1

- How many n-grams are there in a sentence of length m?

Trigram probabilities of a sentence



$$P_{un}(I 'm sorry, Dave. (/s)) = 2.83 \times 10^{-9}$$

$$P_{bi}(I 'm sorry, Dave. (/s)) = 0.33$$

$$P_{tri}(I 'm sorry, Dave. (/s)) = 0.50$$

Short detour: colorless green ideas

But it must be recognized that the notion 'probability of a sentence' is an entirely useless one, under any known interpretation of this term. — Chomsky (1968)

- The following 'sentences' are categorically different:
 - Furiously sleep ideas green colorless
 - Colorless green ideas sleep furiously
- Can n-gram models model the difference?
- Should n-gram models model the difference?

What do n-gram models model?

- Some morphosyntax: the bigram 'ideas are' is (much) more likely than 'ideas is'
- Some semantics: 'bright ideas' is more likely than 'green ideas'
- Some cultural aspects of everyday language: 'Chinese food' is more likely than 'British food'
- more aspects of 'usage' of language

How to test n-gram models?

Extrinsic: Improvement of the target application due to the language model:

- Speech recognition accuracy
- BLEU score for machine translation
- Keystroke savings in predictive text applications

Intrinsic: the higher the probability assigned to a test set better the model. A few measures:

- Likelihood
- (cross) entropy
- perplexity

Like any ML method, test set has to be different than training set.

Intrinsic evaluation metrics: cross entropy

- Cross entropy of a language model on a test set w is

$$H(w) = -\frac{1}{N} \sum_{w_i} \log_2 \hat{P}(w_i)$$

- The lower the cross entropy, the better the model
- Cross entropy is not sensitive to the test-set size

Reminder: Cross entropy is the bits required to encode the data coming from P using another (approximate) distribution \hat{P} .

$$H(P, Q) = -\sum_x P(x) \log \hat{P}(x)$$

Intrinsic evaluation metrics: perplexity

- Perplexity is a more common measure for evaluating language models

$$PP(w) = 2^{H(w)} = P(w)^{-\frac{1}{N}} = \sqrt[N]{\frac{1}{P(w)}}$$

- Perplexity is the average branching factor
- Similar to cross entropy
 - lower better
 - not sensitive to test set size

What do we do with unseen n-grams?

...and other issues with MLE estimates

- Words (and word sequences) are distributed according to the Zipf's law: many words are rare.
- MLE will assign 0 probabilities to unseen words, and sequences containing unseen words
- Even with non-zero probabilities, MLE overfits the training data
- One solution is **smoothing**: take some probability mass from known words, and assign it to unknown words



Laplace smoothing

(Add-one smoothing)

- The idea (from 1790): add one to all counts
- The probability of a word is estimated by

$$P_{+1}(w) = \frac{C(w)+1}{N+V}$$

N : number of word **tokens**
 V : number of word **types** - the size of the vocabulary

- Then, probability of an unknown word is:

$$\frac{0+1}{N+V}$$

Laplace smoothing

for n-grams

- The probability of a bigram becomes

$$P_{+1}(w_1 w_2) = \frac{C(w_1 w_2) + 1}{N + V^2}$$

- and, the conditional probability

$$P_{+1}(w_2 | w_1) = \frac{C(w_1 w_2) + 1}{C(w_1) + V}$$

- In general

$$P_{+1}(w_1^{i-1} w_i) = \frac{C(w_1^{i-1} w_i) + 1}{N + V^n}$$

$$P_{+1}(w_1^{i-1} w_i | w_1^{i-2} w_{i-1}) = \frac{C(w_1^{i-2} w_{i-1} w_i) + 1}{C(w_1^{i-2} w_{i-1}) + V}$$

Bigram probabilities

MLE vs. Laplace smoothing

$w_1 w_2$	C_{+1}	$P_{MLE}(w_1 w_2)$	$P_{+1}(w_1 w_2)$	$P_{MLE}(w_2 w_1)$	$P_{+1}(w_2 w_1)$
(s) I	3	0.118	0.019	1.000	0.188
I 'm	3	0.118	0.019	0.667	0.176
'm sorry	2	0.059	0.012	0.500	0.125
sorry, /	2	0.059	0.012	1.000	0.133
, Dave	2	0.059	0.012	1.000	0.133
Dave. /	2	0.059	0.012	1.000	0.133
I 'm afraid	2	0.059	0.012	0.500	0.125
afraid I	2	0.059	0.012	1.000	0.133
I can	2	0.059	0.012	0.333	0.118
can't	2	0.059	0.012	1.000	0.133
n't do	2	0.059	0.012	1.000	0.133
do that	2	0.059	0.012	1.000	0.133
that.	2	0.059	0.012	1.000	0.133
./ (/s)	3	0.118	0.019	1.000	0.188
Σ		1.000	0.193		

MLE vs. Laplace probabilities

probabilities of sentences and non-sentences (based on the bigram model)

w	I	'm	sorry	,	Dave	.	(/s)	
P_{MLE}	1.00	0.67	0.50	1.00	1.00	1.00	0.33	
P_{+1}	0.19	0.18	0.13	0.13	0.13	0.13	0.19	1.84×10^{-6}

w	,	'm	I	.	sorry	Dave	(/s)	
P_{MLE}	0.00	0.00	0.00	0.00	0.00	0.00	0.00	
P_{+1}	0.03	0.03	0.03	0.03	0.03	0.03	0.03	1.17×10^{-12}

w	I	'm	afraid	,	Dave	.	(/s)	
P_{MLE}	1.00	0.67	0.50	0.00	1.00	1.00	0.00	
P_{+1}	0.19	0.18	0.13	0.03	0.13	0.13	0.19	4.45×10^{-7}

Lidstone correction

(Add- α smoothing)

- A simple improvement over Laplace smoothing is adding α instead of 1

$$P_{+\alpha}(w_{i-n+1}^i | w_{i-n+1}^{i-1}) = \frac{C(w_{i-n+1}^i) + \alpha}{C(w_{i-n+1}^i) + \alpha V}$$

- With smaller α values, the model behaves similar to MLE, it overfits: it has high variance
- Larger α values reduce overfitting/variance, but result in large bias

We need to tune α like any other hyperparameter.

Good-Turing smoothing

- Estimate the probability mass to be reserved for the novel n-grams using the observed n-grams
- Novel events in our training set is the ones that occur once

$$p_0 = \frac{n_1}{n}$$

where n_1 is the number of distinct n-grams with frequency 1 in the training data

- Now we need to discount this mass from the higher counts
- The probability of an n-gram that occurred r times in the corpus is

$$(r+1) \frac{n_{r+1}}{n_r n}$$

Issues with Good-Turing discounting

With some solutions

- Zero counts: we cannot assign probabilities if $n_{r+1} = 0$
- The estimates of some of the frequencies of frequencies are unreliable
- A solution is to replace n_r with smoothed counts z_r
- A well-known technique (simple Good-Turing) for smoothing n_r is to use linear interpolation

$$\log z_r = a + b \log r$$

Back-off and interpolation

The general idea is to fall-back to lower order n-gram when estimation is unreliable

- Even if, $C(\text{black squirrel}) = C(\text{black wug}) = 0$

it is unlikely that

$$C(\text{squirrel}) = C(\text{wug})$$

in a reasonably sized corpus

Interpolation

Interpolation uses a linear combination:

$$P_{int}(w_i | w_{i-1}) = \lambda P(w_i | w_{i-1}) + (1 - \lambda) P(w_i)$$


In general (recursive definition),

$$P_{int}(w_i | w_{i-1}^{i-1}) = \lambda P(w_i | w_{i-1}^{i-1}) + (1 - \lambda) P_{int}(w_i | w_{i-1}^{i-2})$$

- $\sum \lambda_i = 1$
- Recursion terminates with
 - either smoothed unigram counts
 - or uniform distribution $\frac{1}{V}$

How much probability mass does +1 smoothing steal?

- Laplace smoothing reserves probability mass proportional to the size of the vocabulary
- This is just too much for large vocabularies and higher order n-grams
- Note that only very few of the higher level n-grams (e.g., trigrams) are possible




Absolute discounting



- An alternative to the additive smoothing is to reserve an explicit amount of probability mass, c , for the unseen events
- The probabilities of known events has to be re-normalized
- How do we decide what c value to use?

Good-Turing example



$$P_{GT}(\text{the}) + P_{GT}(\text{a}) + \dots = \frac{8}{15}$$

$$P_{GT}(\text{that}) = P_{GT}(\text{do}) = \dots = \frac{2 \times 2}{15 \times 8}$$

$$P_{GT}(\text{'m}) = P_{GT}(\text{'}) = \frac{3 \times 1}{15 \times 2}$$

Not all (unknown) n-grams are equal



- Let's assume that black squirrel is an unknown bigram
- How do we calculate the smoothed probability

$$P_{+1}(\text{squirrel} | \text{black}) = \frac{0 + 1}{C(\text{black}) + V}$$

- How about black wug?

$$P_{+1}(\text{black wug}) = P_{+1}(\text{wug} | \text{black}) = \frac{0 + 1}{C(\text{black}) + V}$$

- Would it make a difference if we used a better smoothing method (e.g., Good-Turing?)

Back-off

Back-off uses the estimate if it is available, 'backs off' to the lower order n-gram(s) otherwise:

$$P(w_i | w_{i-1}) = \begin{cases} P^*(w_i | w_{i-1}) & \text{if } C(w_{i-1} w_i) > 0 \\ \alpha P(w_i) & \text{otherwise} \end{cases}$$

where,

- $P^*(\cdot)$ is the discounted probability
- α makes sure that $\sum P(w)$ is the discounted amount
- $P(w_i)$, typically, smoothed unigram probability

Some shortcomings of the n-gram language models

The n-gram language models are simple and successful, but ...

- They cannot handle long-distance dependencies: In the last race, the horse he bought last year finally _____.
- The success often drops in morphologically complex languages
- The smoothing methods are often 'a bag of tricks'
- They are highly sensitive to the training data: you do not want to use an n-gram model trained on business news for medical texts

Cluster-based n-grams

- The idea is to cluster the words, and fall-back (back-off or interpolate) to the cluster
 - For example,
 - a clustering algorithm is likely to form a cluster containing words for food, e.g., {apple, pear, broccoli, spinach}
 - if you have never seen eat your broccoli, estimate
- $$P(\text{broccoli} | \text{eat your}) \approx P(\text{FOOD} | \text{eat your}) \times P(\text{broccoli} | \text{FOOD})$$
- Clustering can be
 - hard a word belongs to only one cluster (simplifies the model)
 - soft words can be assigned to clusters probabilistically (more flexible)

Modeling sentence types

- Another way to improve a language model is to condition on the sentence types
- The idea is different types of sentences (e.g., ones related to different topics) have different behavior
- Sentence types are typically based on clustering
- We create multiple language models, one for each sentence type
- Often a 'general' language model is used, as a fall-back

Structured language models

- Another possibility is using a generative parser
- Parsers try to explicitly model (good) sentences
- Parsers naturally capture long-distance dependencies
- Parsers require much more computational resources than the n-gram models
- The improvements are often small (if any)

Neural language models

- Similar to maxent models, we can train a feed-forward network that predicts a word from its context
- (gated) Recurrent networks are more suitable to the task:
 - Train a recurrent network to predict the next word in the sentence
 - The hidden representations reflect what is useful in the history
- Combined with embeddings, RNN language models are generally more successful than n-gram models
- In recent years, *masked language models*, combined with neural network architectures called Transformers became the dominant language models

Additional reading, references, credits

- Textbook reference: Jurafsky and Martin (2009, chapter 4) (draft chapter for the 3rd version is also available). Some of the examples in the slides come from this book.
- Chen and J. Goodman (1998) and Chen and J. Goodman (1999) include a detailed comparison of smoothing methods. The former (technical report) also includes a tutorial introduction
- J. T. Goodman (2001) studies a number of improvements to (n-gram) language models we have discussed. This technical report also includes some introductory material
- Gale and Sampson (1995) introduce the 'simple' Good-Turing estimation noted on Slide 12. The article also includes an introduction to the basic method.

Skipping

- The contexts
 - boring (the lecture was
 - boring (the) lecture yesterday was
- are completely different for an n-gram model
- A potential solution is to consider contexts with gaps, 'skipping' one or more words
- We would, for example model $P(e | a b c d)$ with a combination (e.g., interpolation) of
 - $P(e | a b c _)$
 - $P(e | a b _ d)$
 - $P(e | a _ c d)$
 - ...

Caching

- If a word is used in a document, its probability of being used again is high
- Caching models condition the probability of a word, to a larger context (besides the immediate history), such as
 - the words in the document (if document boundaries are marked)
 - a fixed window around the word

Maximum entropy models

- We can fit a logistic regression 'max-ent' model predicting $P(w | \text{context})$
- Main advantage is to be able to condition on arbitrary features

Summary

- We want to assign probabilities to sentences
- N-gram language models do this by
 - estimating probabilities of parts of the sentence (n-grams)
 - use the n-gram probability and a conditional independence assumption to estimate the probability of the sentence
- MLE estimate for n-gram overfit
- Smoothing is a way to fight overfitting
- Back-off and interpolation yields better 'smoothing'
- There are other ways to improve n-gram models, and language models without (explicitly) use of n-grams

Next:

- (?) Neural language models
- Tokenization / Computational morphology

Additional reading, references, credits (cont.)

- The quote from 2001: A Space Odyssey, 'I'm sorry Dave. I'm afraid I can't do it.' is probably one of the most frequent popular culture quotes in the CL literature. It was also quoted, among many others, by Jurafsky and Martin (2009).
- The HAL9000 camera image on page 12 is from Wikipedia, (re)drawn by Wikipedia user Cryteria.

- [1] Chen, Stanley P and Jeffrey Goodman (1998). *An empirical study of smoothing techniques for language modeling*. Tech. rep. TR-98-06, Stanford University Computer Science Group. <http://nlp.stanford.edu/IRN/html/98/06/06.html>
- [2] — (1999). *An empirical study of smoothing techniques for language modeling*. In *Computer speech & language* 13.4, pp. 389–396.
- [3] Choumou Ntuen (1998). *Quora's empirical assumptions*. In *Spilhaus* 19.1, pp. 55–66. [doi:10.1007/BF00000000](https://doi.org/10.1007/BF00000000)
- [4] Gale, William A and Geoffrey Sampson (1995). *Good-Turing frequency estimation without tears*. In *Journal of Quantitative Linguistics* 2.5, pp. 217–227.
- [5] Goodhew, Joshua T (2001). *A bit of progress in language modeling without tears*. Tech. rep. 5000.70.2001.70, Microsoft Research.
- [6] Goodhew, Joshua and James H. Martin (2005). *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition*. second. Pearson Education (Add. com. 0-0-0-01-00000-0)

Additional reading, references, credits (cont.)

- [7] Reddy, Richard (1998). *Lexical Hypotheses in Continuous Speech*. In *Cognitive Models of Speech Processing*, Ed. by Gary T. M. Althaus. MIT Press.

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.2

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.3

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.7

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.8

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.9

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.10

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.11

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.12

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.13

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.14

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.15

C. Gökdeniz

SR | University of Tübingen

November/December 2021

A.16